

# Pandas Notes

Jihuan Zhang

## 1 Reading/Saving Datasets

### 1.1 Reading Datasets

```
data = pd.read_csv("../parent_folder/file.csv", # '_excel' for .xls; '_table' for .txt; '_json' for .json
header=None, # 1st row is data; =0: 1st row is column name
nrows=20, # only read first 20 rows
skiprows=0, # skip nothing; =20: skip first 20 rows; =range(0,21,2); =lambda x: x%2==1: skip odd rows
usecols=[0,1,2], # only read columns 0,1,2; can also be a list of column names
index_col=['positionId'], # use specific column for index; ow a new column will be created for index
keep_default_na=True, # set missing values as NaN; =False: set missing values as blank
na_values=[0,'NA'], # set all zero and 'NA' as NaN
na_filter=False, # do not do anything to the missing values, display the original values
dtype={'positionId':str, 'companyName':str}, # customize data type
parse_dates=['Year'] # set column 'Year' to be time series data
)
```

### 1.2 Creating Datasets

```
data = [[1,2,3],[4,5,6]] # each element represent a row in dataframe
pd.DataFrame(data) # ndarray (structured or homogeneous), Iterable, dict, or DataFrame
pd.DataFrame(data, index=['row1','row2'], columns=['col1','col2','col3'])
pd.DataFrame(data, index=range(len(data)))
data = [1,2,3,4,5]
pd.DataFrame(Counter(data).elements()) # dict
pd.DataFrame(Counter(data).values()) # dict
pd.DataFrame(range(5)) # Iterable
pd.DataFrame(permutations(data)) # Iterable
```

### 1.3 Saving Datasets

```
df.to_csv("out.csv", # output file name
encoding='utf-8',
index=False, # a new column will be created for index
na_rep='missing' # customize NaN
)
```

## 2 Attributes and Methods

Attributes have no parentheses, e.g., `df.shape`, `df.values`. Methods have parentheses, e.g., `df.sum()`, `df.dropna()`

## 2.1 Descriptive Statistics

```
df['A'].sum()           # Sum of all values
df['A'].mean()          # Average
df['A'].median()        # Median value
df['A'].min()           # Minimum value
df['A'].max()           # Maximum value
df['A'].std()           # Standard deviation
df['A'].var()           # Variance
df['A'].count()         # Count non-null values
df['A'].nunique()        # Count unique values
df['A'].quantile(0.75)  # 75th percentile
df['A'].prod()          # Product of all values
```

## 2.2 Data Inspection

```
df.head(n)             # First n rows (default 5)
df.tail(n)             # Last n rows (default 5)
df.shape               # (rows, columns)
df.info()              # Data types and memory usage
df.describe()          # Statistical summary
df.columns             # Column names
df.dtypes              # Data types of each column
df.index              # Index information
df.values              # Numpy array of values
df.memory_usage()      # Memory usage per column
```

# 3 Displaying Datasets

## 3.1 Summary Descriptive Statistics

```
df.shape[0] # df.shape is a tuple. df.shape[0] is number of rows, df.shape[1] is number of columns
len(df) # this actually works and gives you the number of rows
df.describe() # descriptive statistics only for numeric-type columns
df.describe(include='O') # descriptive statistics only for object-type columns (strings, list, dict, mixed); count
                        : Number of non-null values; unique: Number of unique values; top: Most frequently occurring value; freq:
                        : Frequency of the most common value
df.describe(include='all')
df.describe().round(4) # round the results in 4 decimal spaces
df.describe().round(4).T # transpose
```

## 3.2 Displaying sub-dataframe by Indexing

```
df[['col1', 'col2']] # display a sub-dataframe of two columns
df[['col1']] # display a sub-dataframe of one column
df['col1'] # this returns a Series instead of a dataframe; a Series in Pandas has attributes 'name' (instead of '
           : columns'), 'index' and 'values'; Series has no 'columns' attribute; df['col1']+*//df['col2'] is
           : elementwise, and the result is still a Series; for calculations with vector, e.g, inner and outer product, we
           : need to use 'df.values' which is an array, e.g., df['col1'].values @ df['col2'].values

df.iloc[:, [0]] # integer-location based indexing; df.iloc() is wrong
df.iloc[:, 0] # similarly, this returns a Series instead of a dataframe
```

```

df.iloc[0,:] # this is also a Series! df.iloc[0,:].name will be the corresponding index of df; df.iloc[0,:].index
            will be the corresponding columns of df
df.iloc[0,0] # this returns a scalar

df.loc[index_1:index_n, 'col_1':'col_n'] # location (label) based indexing; if the index is labeled by 0,1,2,...,
    then df.loc[1:3,:] works and returns row 1 to 3 (INCLUDING 3); if the index is labeled differently, say 'a','b',
    'c',..., then df.loc[1:3,:] will return an error
df.loc[index_1,:] # similarly, this returns a Series instead of a dataframe
df.loc[index_1, 'col_1'] # similarly, this returns a scalar

series.to_frame() # convert Series to DataFrame
pd.DataFrame(series) # convert Series to DataFrame, more general
pd.DataFrame([scalar]) # convert scalar to DataFrame, more general

```

### 3.3 Displaying sub-dataframe by Filtering Conditions

```

df[df['salary']<=70000].iloc[:,1:3] # if you use R, you might have intuitively tried df.iloc[df['salary']<=70000],1:3]. This is wrong in pandas
df[df['name'].isin(['Joe','Henry','Delta'])] # df['name'] in (['Joe','Henry','Delta']) is wrong

df['col1'].iloc[0] # a scalar; since df['col1'] is a Series, iloc only has one dimension; same as df[['col1']].iloc[0,0]

# for Series, operator must be &, | instead of and, or, and each condition must be in parenthesis ()
df[(df['col1'] < 1000) | (df['col2'].isna())] # df[df['col1'] < 1000 or df['col2'].isna()] is wrong

df[df['col1'].between(val1, val2)] # between

# filtering strings
df['col1'].str.startswith('abc')
df['col1'].str.endswith('abc')
df['col1'].str.contains('abc')

df.iloc[:, [i%2==1 for i in range(len(df.columns))]]
df.iloc[:, [j for j in range(1, len(df.columns)+1, 2)]]
df.loc[:, df.columns.str.endswith('数')]

```

## 4 Modifying Dataframes

### 4.1 Changing Column Names

```

df.rename(columns={'old_name': 'new_name'}) # change a small number of column names

# changing a large number of column names
# f"..." is a formatted string literal (an f-string). It lets you put expressions inside "..." that are evaluated
    and converted to text.
# enumerate(df.columns) convert column names to Iterable, whose elements are [0, 'colname1'], [1, 'colname2'], ...
df.columns = [f"X_{i+1}" if i <= 100 else c for i, c in enumerate(df.columns)]

# equivalently
n = 2 # 要改的列数: 第0~(2-1)列
new_names = [f"X_{i}" for i in range(1, n+1)]
df.rename(columns=dict(zip(df.columns[:n], new_names)), inplace=True)

# equivalently

```

```
n = 2
new_names = list(df.columns)
for i in range(0,n):
    new_names[i] = f'X_{i+1}'
df.columns = new_names
```

## 4.2 Modifying Values

```
# string values
Series.str.title() # Converts first character of each word to uppercase and remaining to lowercase.
Series.str.capitalize() # Converts first character to uppercase and remaining to lowercase.
Series.str.swapcase() # Converts uppercase to lowercase and lowercase to uppercase.
```

## 4.3 Sort

```
df.sort_values(by, # str or list of str; Name or list of names to sort by. if axis is 0 or 'index' then by may
               contain index levels and/or column labels. if axis is 1 or 'columns' then by may contain column levels and/
               or index labels.
               axis=0, # "{0 or 'index' , 1 or 'columns' }", default 0. Axis to be sorted.
               ascending=True, # bool or list of bool, default True. Sort ascending vs. descending. Specify list for multiple
               sort orders. If this is a list of bools, must match the length of the by.
               inplace=False # bool, default False. If True, perform operation in-place.
               na_position='last', # { 'first' , 'last' }, default 'last' . Puts NaNs at the beginning if first; last puts
               NaNs at the end.
               )
```

## 4.4 Adding New Columns

### 4.4.1 Assign

`assign()` is a method for adding new columns to a DataFrame in a functional, chainable way. Unlike direct assignment (`df['col'] = value`), `assign()` returns a new DataFrame without modifying the original.

```
df.assign(new_col1 = value1, new_col2 = value2, ...) # create new columns with identical values
df3 = df.assign(new_col1=df['col1']*12, new_col2=df['new_col1']/1000, new_col3=df['new_col2']>60000, ...) # you
               can create new columns using other new columns

df.assign(
    bonus=lambda x: x['salary'] * 0.1,
    total_comp=lambda x: x['salary'] + x['bonus'], # Uses 'bonus' created above!
    tax=lambda x: x['total_comp'] * 0.25
) # Creating New Columns Based on Anonymous Functions
```

### 4.4.2 Advanced Methods

```
# Custom numerical Grouping with pd.cut()
df['age_group'] = pd.cut(df['age'], bins=[0, 10, 20, 30], right=True, include_lowest=True) # right side inclusive;
               include_lowest=True: Without include_lowest, 0 would be NaN
# Instead of equal-width bins, create equal-frequency bins (quantiles)
df['age_quartile'] = pd.qcut(df['age'], q=4, labels=['Q1', 'Q2', 'Q3', 'Q4'])

# Custom string Grouping with map()
```

```

rating_groups = {
    'Good': 'High',
    'Excellent': 'High',
    'Normal': 'Low',
    'Bad': 'Low'
}
df['rating_group'] = df['rating'].map(rating_groups) # Create new column with groups

```

## 4.5 Dropping NAs/Duplicates

```

df.dropna(axis=0, # {0 or 'index', 1 or 'columns'}, default 0; =0 or 'index': drop rows which contain missing
               values; =1 or 'columns': drop columns which contain missing values; only a single axis is allowed.
          how='any', # {'any', 'all'}, default 'any'
          subset=None, # default None; Labels along other axis to consider, e.g. if you are dropping rows these would be a
                        list of columns to include.
          inplace=False, # bool, default False; Whether to modify the DataFrame rather than creating a new one.
        )
df.dropna() # return a dataframe that drops all rows that contain at least one NA

df.drop_duplicates(subset=None,
                  keep='first', # {'first', 'last', False}, default 'first' Determines which duplicates (if any) to keep. '
                                first' : Drop duplicates except for the first occurrence. 'last' : Drop duplicates except for the last
                                occurrence. False : Drop all duplicates.
                  inplace=False # bool, default False; Whether to modify the DataFrame rather than creating a new one.
        )
df.drop_duplicates() # drop all duplicates except for the first occurrence; a duplicate is define by having the
                    same row across all columns

```

## 4.6 Displaying NAs/Duplicates (Advanced)

```

df.isna() # return a dataframe where missing values are False and others are True
df.isna().sum() # return a Series of count of missing values per column
df.isna().sum().sum() # return a scalar of total count of missing values

df.isna().any() # per column across rows: return a Series with unnamed name and indices given by columns names;
                False means the corresponding COLUMN contains no NA
df.isna().any(axis=1) # per row across columns: return a Series with unnamed name and indices given by the indices
                      of df; False means the corresponding ROW contains no NA
df.isna().all() # per column across rows: return a Series with unnamed name and indices given by columns names;
                True means all rows in the corresponding COLUMN are NAs

df[df.isna().any(axis=1)] # return a dataframe containing rows with NA; axis=1 means by row, default is by column;
                          of course, df[df.isna().any()] would not work
df[df[['col1', 'col2']].isna().any(axis=1)] # return a dataframe whose 'col1' or 'col2' contains rows with NA; axis
=1 means by row, default is by column
df[df.isna().sum(axis=1) > N] # Rows with more than N NAs
df.isna().sum(axis=1) # Count NAs per row
df.columns[df.isna().any()] # column names that contain NAs

df.fillna('*') # return a dataframe whose missing values are set to be '*'

df.duplicated() # return a series of bools, True mean this row is a duplicate
df[df.duplicated()] # return duplicated rows
df[df.duplicated(['col1', 'col2'])] # find rows whose values in 'col1' and 'col2' as a tuple, is duplicated

```

## 4.7 Ranking

```
# return a Series
df['new_col'] = df['col_to_be_ranked'].rank(axis=0, method='average', numeric_only=False, na_option='keep',
      ascending=True, pct=False) # method:{ 'average' , 'min' , 'max' , 'first' , 'dense' }, default 'average'
```

## 5 Groupby

```
# df.groupby(...) returns a GroupBy object, which is different from DataFrame and Series
df.groupby(['col1','col2'])['col3'].agg('min') # group by ('col1','col2') and calculate min for 'col3' for each
      group; this returns a Series with name 'col3' and index ['col1','col2']
df.groupby(['col1','col2'])['col3'].agg('min').reset_index() # this reset the index, changing ['col1','col2'] to
      columns, which make the previous Series to a DataFrame
df.groupby(['col1','col2'])[['col3']].agg('min') # note that [['col3']] makes the result a DataFrame instead of a
      Series.
df.groupby(['col1','col2'])['col3'].agg(['min']) # this changes the name 'col3' into 'min', so it is slightly
      different from .agg(['min'])
df.groupby(['col1','col2'])['col3'].min() # indeed, if we only want one descriptive statistic, we don't have to use
      .agg() method

df.groupby('category')['value'].agg([
    'count',      # Count of non-null values
    'sum',        # Sum of values
    'mean',       # Average (mean)
    'median',     # Median value
    'min',        # Minimum value
    'max',        # Maximum value
    'std',        # Standard deviation
    'var',        # Variance
    'first',      # First value
    'last',       # Last value
    'nunique',    # Number of unique values
])

df.groupby('col1').agg(new_colname1 = ('col1','min'), new_colname2 = ('col2','max'), ...)
df.groupby('col1')['col2'].agg([('new_col1','min'), ('new_col2','max'), ...])

# one can customize functions in agg()
df.groupby('category')['value'].agg(null_count=lambda x: x.isna().sum()) # here x represent each group, which is a
      Series in this case (['value'] instead of [['value']]); x.isna() is the same class as x, where every entry is
      set True for NA and False otherwise; sum()=sum(axis=0) by default, per column across rows

sales.groupby('region')['revenue'].agg([
    'mean',
    ('25th_percentile', lambda x: x.quantile(0.25)),
    ('median', 'median'),
    ('75th_percentile', lambda x: x.quantile(0.75))
])
```

## 6 Merge

```
pd.merge(left, right, how='inner', on=['key1', 'key2'])
pd.merge(left, right, on=['key1', 'key2']) # how='inner' is default
pd.merge(left, right, how='outer', on=['key1', 'key2'])
```

```
pd.merge(left, right, how='right', on=['key1', 'key2'])
pd.merge(left, right, on='k', suffixes=['_l', '_r'])
left.merge(right, how='inner', on=None, left_on=None, right_on=None, left_index=False, right_index=False, sort=False, suffixes=('_x', '_y'), copy=None, indicator=False, validate=None)
```

## 7 Pivit Table

```
df.pivot_table(values=None, # Column(s) to aggregate
               index=None, # Column(s) for row labels
               columns=None, # Column(s) for column labels
               aggfunc='mean', # default "mean"
               fill_value=None, # Value to replace NaN
               margins=False, # If margins=True, special All columns and rows will be added with partial group aggregates
                             # across the categories on the rows and columns.
               dropna=True # If True, rows with a NaN value in any column will be omitted before computing margins
               )
```

### 7.1 Naming Rule for Columns by Pivit Tables

```
# Level 0: one of the aggfunc (if more than more)
# Level 1: one of the values
# Level 2: columns[0]
# Level 3: columns[1]
# Level 4: columns[2] (if you have 3 columns parameters)
# ...

# When: single values + single columns + single aggfunc, result has simple column names
# ['product_val_1', 'product_val_2', ...]

# When: multiple values + multiple columns + single aggfun, result is a tuple
# [('value_name_1', 'col1_val1', 'col2_val1', ...), ..., ('value_name_2', 'col1_val1', 'col2_val1', ...), ...]

# When: multiple values + multiple columns + multiple aggfun, result is a tuple
# [('aggfunc_1', 'value_name_1', 'col1_val1', 'col2_val1', ...), ..., ('aggfunc_2', 'value_name_1', 'col1_val1', 'col2_val1', ...), ...]
```

## 8 Time Data

### 8.1 Creating Time Data

```
date = pd.to_datetime('2015-01-25')
```

### 8.2 Displaying Time Information

```
date.day_name() # returns 'Sunday'
date.weekday() # returns 6; Monday is 0, Sunday is 6
date.year
date.quarter
date.isocalendar().week # i-th week of the year
date.day # i-th day of the month
```

```
date.month  
date.month_name()
```

## 8.3 Converting Time Unit

```
df['col1'].dt.total_seconds() # convert the time into seconds
```

## 8.4 Calculating Times

```
df['time_col'].iloc[1] - df['time_col'].iloc[0] == pd.Timedelta(days=1) # Available kwargs: {days, seconds, microseconds, milliseconds, minutes, hours, weeks}. Values for construction in compat with datetime.timedelta. Numpy ints and floats will be coerced to python ints and floats. years/months does not work because Timedelta represents a fixed duration, so it can't handle variable-length periods  
  
pd.to_datetime('2015-01-30') + pd.DateOffset(months=1) # returns Timestamp('2015-02-28 00:00:00'), automatically offset backwards; Available kwargs: {years, months, days, seconds, microseconds, milliseconds, nanoseconds, minutes, hours, weeks}.
```

## 8.5 Grouper

```
df.groupby(pd.Grouper(key='date',freq='ME'))
```

# 9 Other Functions/Methods

```
isinstance(n, int) and n > 0 # check if n is a natural number  
f'row_{i}' # makes i changeable
```