

MySQL Notes

Jihuan Zhang

1 Basic Syntax

1.1 Displaying Subtables

- The **SELECT** ... **FROM** ... statement is used to select and display data from a database.

```
SELECT * FROM table1; -- select the entire table
SELECT col1, col2, ... FROM table1; -- select all rows for specific columns from a table
```

- The **WHERE** clause is used to filter rows. Some basic logic operators are: **AND**, **OR**, **NOT**, **=** (not “==”), **!=**/**<>**, **<**, **<=**, **>**, **>=**.

```
SELECT * FROM table1 WHERE col1 <= num AND col2 = 'string1';
SELECT * FROM table1 WHERE col1 <= col2; -- operators in sql are vectorized
```

- ORDER BY** is used to order the result. Default is ascending order. Use **ASC**/**DESC** to specify order.

```
SELECT * FROM table1 ORDER BY col1; -- also works for strings in alphabetical order
SELECT * FROM table1 ORDER BY col1 ASC; -- same as above
SELECT * FROM table1 ORDER BY col1 DESC; -- specify descending order
SELECT * FROM table1 ORDER BY col1, col2; -- order wrt col1, and within the same values of col1, order wrt col2
SELECT * FROM table1 ORDER BY col1 DESC, col2 ASC, ...; -- customize the orders
```

- LIMIT** and **OFFSET** are used to restrict the number of rows to display¹.

```
SELECT * FROM table1 ORDER BY col1 LIMIT 3; -- display the first 3 rows (ordered)
SELECT * FROM table1 ORDER BY col1 DESC LIMIT 3; -- display the last 3 rows
SELECT * FROM table1 ORDER BY col1 LIMIT 3 OFFSET 1; -- display the first 3 rows ignoring row 1 (i.e., rows 2-4)
SELECT * FROM table1 ORDER BY col1 LIMIT 3;
```

- AS** (Aliases) for subqueries and new variable name

```
SELECT col1 AS new_name FROM (SELECT col1, col2 FROM table1 ORDER BY col3) AS new_table_name;
```

1.2 More on Operators

- BETWEEN** num1 **AND** num2: $\text{num1} \leq \cdot \leq \text{num2}$. In many SQL versions, num1 must not be greater than num2.

```
SELECT * FROM table1 WHERE col1 BETWEEN 2*PI() AND SQRT(2500);
```

$\pi = \text{PI}()$, $\sqrt{\cdot} = \text{SQRT}(\cdot)$.

- LIKE** 'string' (case-insensitive); **LIKE BINARY** 'string' (case-sensitive).

Wildcards (通配符): **%** (zero or more arbitrary characters), **_** (one arbitrary character).

```
SELECT * FROM table1 WHERE col1 LIKE 'a%'; -- start with 'a'
SELECT * FROM table1 WHERE col1 LIKE '%a'; -- end with 'a'
SELECT * FROM table1 WHERE col1 LIKE '%a%'; -- contain 'a'
SELECT * FROM table1 WHERE col1 LIKE '_a%' -- second character is 'a'
SELECT * FROM table1 WHERE col1 LIKE '%a_' -- second to last character is 'a'
SELECT * FROM table1 WHERE col1 LIKE '%a_%' -- contain 'a' followed by at least one character
SELECT * FROM table1 WHERE col1 LIKE '%_a_%' -- contain 'a' in between two characters
```

¹They are supported by SQL Server and Oracle.

- A space is a character in SQL.

Operator precedence (运算符优先级): `()` > everything else > `NOT` > `AND` > `OR`.

- `IN()`:

```
SELECT * FROM table1 WHERE col1 IN (num1, num2, num3);
SELECT * FROM table1 WHERE col1 = num1 OR col1 = num2 OR col1 = num3; -- equivalent
SELECT col1 FROM table1 WHERE col2 IN (SELECT col3 FROM table2 WHERE col4 = 'string1'); -- subqueries can be used
```

- `IS NULL`:

```
SELECT * FROM table1 WHERE col1 IS NULL;
SELECT * FROM table1 WHERE col1 IS NULL AND col2 IS NULL;
```

1.3 Modifying a Table

- The `INSERT INTO ... VALUES ...` statement is used to insert new rows in a table.

```
INSERT INTO table1 VALUES (val1, val2, ...); -- insert a row of full length
INSERT INTO table2 (col1, col2, ...) VALUES (val1, val2, ...); -- insert a sub-row
INSERT INTO table1 VALUES (val1, val2, ...), (val1, val2, ...), ...; -- insert multiple rows
```

- The `UPDATE ... SET` statement is used to modify the existing rows.

```
UPDATE table1 SET col1 = val1, col2 = val2, ... WHERE condition;
```

- The `DELETE FROM` statement is used to delete existing rows in a table. The `DROP TABLE` statement delete the entire table.

```
DELETE FROM table1 WHERE condition; -- delete rows satisfying the condition
DELETE FROM table1; -- delete all rows
DROP TABLE table1; -- delete the table
```

- Reverse or commit the modifications using `BEGIN`, `ROLLBACK`, and `COMMIT`.

```
BEGIN; -- Start a transaction/modification

UPDATE accounts SET balance = balance - 100 WHERE id = 1;
UPDATE accounts SET balance = balance + 100 WHERE id = 2;

-- If you change your mind:
ROLLBACK; -- undo both updates

-- If you're sure:
COMMIT; -- make both updates permanent
```

1.4 Aggregate Functions

- `DISTINCT` finds unique values in rows; can be combined with `COUNT()`.

```
SELECT COUNT(*) FROM table1; -- count the number of rows
SELECT DISTINCT col1, col2, ... FROM table1; -- select unique rows wrt (col1, col2, ...)
SELECT COUNT(DISTINCT col1, col2, ...) FROM table1; -- count the number of unique rows wrt (col1, col2, ...)
```

- `COUNT()`, `MIN()`, `MAX()`, `AVG()`, and `SUM()` ignore null values, except for `COUNT(*)`.

```

SELECT MAX(col1) FROM table1 WHERE condition;
SELECT SUM(Quantity * 10) FROM OrderDetails;
SELECT OrderID, SUM(Quantity) AS TotalQuantity FROM OrderDetails GROUP BY OrderID;
SELECT * FROM Products WHERE price > (SELECT AVG(price) FROM Products); -- Return all products with a price
                                above average

```

1.5 Join

- **JOIN ... ON ...**, or **INNER JOIN ... ON ...**, is used to select rows that have matching values in both tables. Use “.” in table1.col1 and table2.col2. **LEFT JOIN**, **RIGHT JOIN**, and **FULL JOIN** are used in the same way.

```

SELECT Orders.OrderID, Customers.CustomerName, Orders.OrderDate
FROM Orders
JOIN Customers ON Orders.CustomerID=Customers.CustomerID;

SELECT Products.ProductID, Products.ProductName, Categories.CategoryName
FROM (Products JOIN Categories ON Products.CategoryID = Categories.CategoryID); -- equivalent, this shows the
                                logic of JOIN

SELECT Orders.OrderID, Customers.CustomerName, Orders.OrderDate
FROM Customers
JOIN Orders ON Orders.CustomerID=Customers.CustomerID; -- equivalent, order of table names doesn't matter

SELECT Orders.OrderID, Customers.CustomerName, Orders.OrderDate
FROM Orders
INNER JOIN Customers ON Orders.CustomerID=Customers.CustomerID; -- equivalent, JOIN is the same as INNER JOIN

```

Join three tables:

```

SELECT Orders.OrderID, Customers.CustomerName, Shippers.ShipperName
FROM ((Orders JOIN Customers ON Orders.CustomerID = Customers.CustomerID)
JOIN Shippers ON Orders.ShipperID = Shippers.ShipperID);

```

- The underlying details of **INNER JOIN**, **LEFT JOIN**, **RIGHT JOIN**, and **FULL JOIN**:

Let the employee table be	ID	Name	DeptID	and the department table be	DeptID	DeptName
	1	Alice	10		10	Sales
	2	Bob	20		20	Marketing
	3	Carol	10		30	IT
	4	David	40		50	Finance

Then, “employee JOIN department ON employee.DeptID = department.DeptID” creates the following table:

employee.ID	employee.Name	employee.DeptID	department.DeptID	DeptID.DeptName
1	Alice	10	10	Sales
2	Bob	20	20	Marketing
3	Carol	10	10	Sales

“employee LEFT JOIN department ON employee.DeptID = department.DeptID” creates the following table:

employee.ID	employee.Name	employee.DeptID	department.DeptID	DeptID.DeptName
1	Alice	10	10	Sales
2	Bob	20	20	Marketing
NULL	NULL	NULL	30	IT
NULL	NULL	NULL	50	Finance

“employee RIGHT JOIN department ON employee.DeptID = department.DeptID” creates the following table:

employee.ID	employee.Name	employee.DeptID	department.DeptID	DeptID.DeptName
1	Alice	10	10	Sales
2	Bob	20	20	Marketing
3	Carol	10	10	Sales
4	David	40	NULL	NULL

“employee FULL JOIN department ON employee.DeptID = department.DeptID” creates the following table:

employee.ID	employee.Name	employee.DeptID	department.DeptID	DeptID.DeptName
1	Alice	10	10	Sales
2	Bob	20	20	Marketing
3	Carol	10	10	Sales
4	David	40	NULL	NULL
NULL	NULL	NULL	30	IT
NULL	NULL	NULL	50	Finance

Therefore, combined with **SELECT ... FROM ...**, it is easy to understand what’s happening. Also, it is clear that table1 LEFT JOIN table2 is the same as table2 RIGHT JOIN table1, with the only difference be the column order.

Now, observe that in the above example, the matching columns have unrepeated values in at least one of the two tables (in the example, Dept.ID in **department** is unrepeated). What if both have repeated values?

When the join columns aren’t unique in either table, you get a Cartesian product of the matching rows - every matching row from table1 is combined with every matching row from table2. For example, let the **department**

table be	DeptID	DeptName	instead.
	10	Sales	
	20	Marketing	
	10	Sales	

Then, “employee JOIN department ON employee.DeptID = department.DeptID” creates the following table:

employee.ID	employee.Name	employee.DeptID	department.DeptID	DeptID.DeptName
1	Alice	10	10	Sales
1	Alice	10	10	Sales
2	Bob	20	20	Marketing
3	Carol	10	10	Sales
3	Carol	10	10	Sales

2 Time Series Data

Create a table consisting of stock name (ticker), time (ts), and prices.

```
CREATE TABLE stock_prices (  
  ticker VARCHAR(10),  
  ts TIMESTAMP,  
  price DECIMAL(10,2)  
);  
  
INSERT INTO stock_prices VALUES  
  ('TSLA', '2025-09-01 09:30:00', 185.20),  
  ('TSLA', '2025-09-01 09:31:00', 185.35),  
  ('TSLA', '2025-09-01 09:32:00', 185.10);  
  
SELECT ts, price,  
  AVG(price) OVER (ORDER BY ts ROWS 4 PRECEDING) AS moving_avg  
FROM stock_prices;
```